

DS-6030 Homework Module 5

Tom Lever

06/24/2023

DS 6030 | Spring 2023 | University of Virginia

8. In this exercise, we will generate simulated data, and will then use this data to perform best subset selection.

- (a) Use the `rnorm()` function to generate a predictor X of length $n = 100$, as well as a noise vector ϵ of length $n = 100$.

```
set.seed(2)
X <- rnorm(n = 100, mean = 0, sd = 1)
epsilon <- rnorm(n = 100, mean = 0, sd = 1*10^{-6})
X[1:3]
```

```
# [1] -0.8969145  0.1848492  1.5878453
```

```
epsilon[1:3]
```

```
# [1]  1.074459e-06  2.605978e-07 -3.142720e-07
```

- (b) Generate a response vector Y of length $n = 100$ according to the model $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \epsilon$, where $\beta_0, \beta_1, \beta_2, \beta_3$ are constants of your choice.

```
beta_0 <- 1
beta_1 <- 2
beta_2 <- 3
beta_3 <- 4
Y <- beta_0 + beta_1 * X + beta_2 * I(X^2) + beta_3 * I(X^3) + epsilon
Y[1:3]
```

```
# [1] -1.266573  1.497471 27.752887
```

- (c) Use the `regsubsets()` function to perform best subset selection in order to choose the best model containing the predictors X, X^2, \dots, X^{10} . What is the best model obtained according to C_p , BIC, and adjusted R^2 ? Show some plots to provide evidence for your answer, and report the coefficients of the best model obtained. Note you will need to use the `data.frame()` function to create a single data set containing both X and Y .


```
data_frame <- data.frame(X = X, Y = Y)
head(data_frame, n = 3)
```



```
#           X           Y
# 1 -0.8969145 -1.26657....
# 2  0.1848492  1.497470....
# 3  1.5878453 27.75288....
```

```

library(leaps)
formula <- Y ~ X + I(X^2) + I(X^3) + I(X^4) + I(X^5) + I(X^6) + I(X^7) + I(X^8) + I(X^9) + I(X
subset_selection_object <- regsubsets(
  x = formula,
  data = data_frame,
  method = "exhaustive"
)
library(TomLeversRPackage)
analyze_subset_selection_object(subset_selection_object)

```





```
# $coefficients_by_Mallows_Cp
# (Intercept)      X      I(X^2)      I(X^3)
#           1         2         3         4
#
```

```
# $coefficients_by_Schwartz_BIC
# (Intercept)      X      I(X^2)      I(X^3)
#           1       2       3       4
#
# $coefficients_by_adjusted_R2
# (Intercept)      X      I(X^2)      I(X^3)
#           1       2       3       4
```


The best model obtained according to Mallows's C_p , the Schwartz Bayesian Information Criterion, and adjusted R^2 is



$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 = 1 + 2X + 3X^2 + 4X^3$$

- (d) Repeat (c), using forward stepwise selection and also using backwards stepwise selection. How does your answer compare to the results in (c)?

The models chosen by forward and backward selection involve the true predictive terms and their coefficients.

```
subset_selection_object <- regsubsets(
  x = formula,
  data = data_frame,
  method = "forward"
)
analyze_subset_selection_object(subset_selection_object)
```





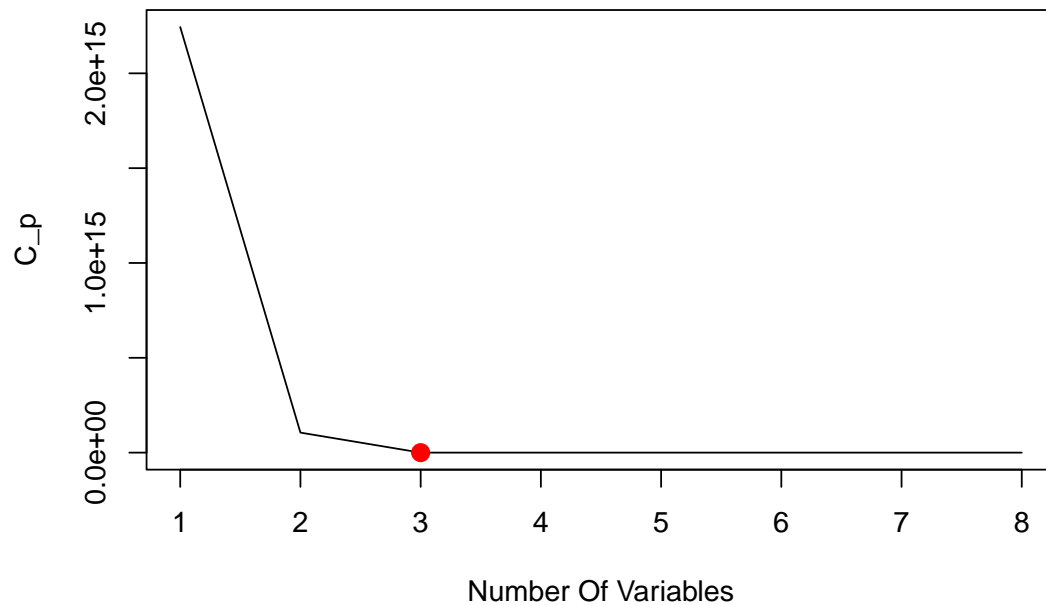
```
# $coefficients_by_Mallows_Cp
# (Intercept)      X      I(X^2)      I(X^3)
#           1         2         3         4
#
```



```

# $coefficients_by_Schwartz_BIC
# (Intercept)      X      I(X^2)      I(X^3)
#           1       2       3       4
#
# $coefficients_by_adjusted_R2
# (Intercept)      X      I(X^2)      I(X^3)
#           1       2       3       4

subset_selection_object <- regsubsets(
  x = formula,
  data = data_frame,
  method = "backward"
)
analyze_subset_selection_object(subset_selection_object)

```





```
# $coefficients_by_Mallows_Cp
# (Intercept)      X      I(X^2)      I(X^3)
#           1         2         3         4
#
```

```
# $coefficients_by_Schwartz_BIC
# (Intercept)      X      I(X^2)      I(X^3)
#           1        2        3        4
#
# $coefficients_by_adjusted_R2
# (Intercept)      X      I(X^2)      I(X^3)
#           1        2        3        4
```


The best model obtained according to Mallows's C_p , the Schwartz Bayesian Information Criterion, and adjusted R^2 is

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 = 1 + 2X + 3X^2 + 4X^3$$

These results are identical to those in part (c).

- (e) Now fit a lasso model to the simulated data, again using X, X^2, \dots, X^{10} as predictors. Use cross-validation to select the optimal value of λ . Create plots of the cross-validation error as a function of λ . Report the resulting coefficient estimates, and discuss the results obtained.

```
full_model_matrix <- model.matrix(object = formula, data = data_frame)[, -1]
the_cv.glmnet <- glmnet::cv.glmnet(x = full_model_matrix, y = Y, alpha = 1)
plot(the_cv.glmnet)
```



```
optimal_value_of_lambda <- the_cv.glmnet$lambda.min
optimal_value_of_lambda
```

```
# [1] 0.3822143
```

```
polynomial_lasso_regression_model <- glmnet::glmnet(full_model_matrix, y = Y, alpha = 1)
predict(object = polynomial_lasso_regression_model, s = optimal_value_of_lambda, type = "coeff
```



```

# 11 x 1 sparse Matrix of class "dgCMatrix"
#               s1
# (Intercept) 1.364119
# X           1.990943
# I(X^2)      2.730444
# I(X^3)      3.899032
# I(X^4)      .
# I(X^5)      .
# I(X^6)      .
# I(X^7)      .
# I(X^8)      .
# I(X^9)      .
# I(X^10)     .

```

The models chosen by polynomial lasso regression involve the true predictive terms and rough approximations of their coefficients. $\beta_0 = 1.364$, $\beta_1 = 1.991$, $\beta_2 = 2.730$, and $\beta_3 = 3.899$.

- (f) Now generate a response vector Y according to the model $Y = \beta_0 + \beta_7 X^7 + \epsilon$, and perform best subset selection and the lasso. Discuss the results obtained.

```

beta_7 <- 8
Y <- beta_0 + beta_7 * I(X^7) + epsilon
data_frame <- data.frame(X = X, Y = Y)
head(data_frame, n = 3)

```

```



#           X           Y
# 1 -0.8969145 -2.73548....
# 2  0.1848492  1.000059....
# 3  1.5878453 204.5857....


```

```

#library(leaps)
formula <- Y ~ X + I(X^2) + I(X^3) + I(X^4) + I(X^5) + I(X^6) + I(X^7) + I(X^8) + I(X^9) + I(X
subset_selection_object <- regsubsets(
  x = formula,
  data = data_frame,
  method = "exhaustive"
)
analyze_subset_selection_object(subset_selection_object)

```






```
# $coefficients_by_Mallows_Cp
# (Intercept)      I(X^7)
#           1           8
#
# $coefficients_by_Schwartz_BIC
# (Intercept)      I(X^7)
#           1           8
#
# $coefficients_by_adjusted_R2
# (Intercept)      I(X^7)
#           1           8
```

The best model obtained according to Mallows's C_p , the Schwartz Bayesian Information Criterion, and adjusted R^2 is

$$Y = \beta_0 + \beta_7 X^7 = 1 + 8X^7$$

```
full_model_matrix <- model.matrix(object = formula, data = data_frame)[, -1]
the_cv.glmnet <- glmnet::cv.glmnet(x = full_model_matrix, y = Y, alpha = 1)
plot(the_cv.glmnet)
```



```
optimal_value_of_lambda <- the_cv.glmnet$lambda.min
optimal_value_of_lambda
```

```
# [1] 19.12382
```

```
polynomial_lasso_regression_model <- glmnet::glmnet(full_model_matrix, y = Y, alpha = 1)
predict(object = polynomial_lasso_regression_model, s = optimal_value_of_lambda, type = "coeff
```

```
# 11 x 1 sparse Matrix of class "dgCMatrix"
```

```
#              s1
# (Intercept) 0.1238844
# X              .
# I(X^2)         .
# I(X^3)         .
# I(X^4)         .
# I(X^5)         .
# I(X^6)         .
# I(X^7)        7.7667958
# I(X^8)         .
# I(X^9)         .
# I(X^10)        .
```

The models chosen by polynomial lasso regression involve the true predictive terms and rough approximations of their coefficients. $\beta_0 = 1.364$ and $\beta_7 = 7.767$.

9. In this exercise, we will predict the number of applications received using the other variables in the College data set.

- (a) Split the data set into a training set and a test set.

```
colleges <- ISLR2::College
head(colleges, n = 3)
```

```
#               Private Apps Accept Enroll Top10perc Top25perc
# Abilene Christian University    Yes 1660   1232    721      23      52
# Adelphi University             Yes 2186   1924    512      16      29
# Adrian College                Yes 1428   1097    336      22      50
#               F.Undergrad P.Undergrad Outstate Room.Board Books
# Abilene Christian University    2885           537    7440    3300    450
# Adelphi University             2683           1227   12280    6450    750
# Adrian College                1036            99   11250    3750    400
#               Personal PhD Terminal S.F.Ratio perc.alumni Expend
# Abilene Christian University    2200   70      78    18.1      12    7041
# Adelphi University             1500   29      30    12.2      16   10527
# Adrian College                1165   53      66    12.9      30    8735
#               Grad.Rate
# Abilene Christian University     60
# Adelphi University              56
# Adrian College                  54
```

```
proportion_of_training_data <- 0.9
library(TomLeversRPackage)
split_data <- split_data_set_into_training_and_testing_data(
  data_frame = colleges,
  proportion_of_training_data = proportion_of_training_data
)
vector_of_indices_of_training_data <- split_data$vector_of_indices_of_training_data
vector_of_indices_of_training_data[1:3]
```

```
# [1] 464 37 678
```

```
vector_of_indices_of_testing_data <- split_data$vector_of_indices_of_testing_data
vector_of_indices_of_testing_data[1:3]
```

```
# [1] 728 491 550
```

```
training_data <- split_data$training_data
head(training_data, n = 3)
```

```
#               Private Apps Accept Enroll Top10perc
# Quincy University    Yes 1025   707    297      22
# Bard College         Yes 1910   838    285      50
# University of Southern California    Yes 12229  8498  2477      45
#               Top25perc F.Undergrad P.Undergrad Outstate
# Quincy University     66           1070      72    10100
# Bard College          85           1004      15    19264
```

# University of Southern California	71	13259	1429	17230
#	Room.Board	Books	Personal	PhD Terminal
# Quincy University	4140	450	1080	69 71
# Bard College	6206	750	750	98 98
# University of Southern California	6482	600	2210	90 94
#	S.F.Ratio	perc.alumni	Expend	Grad.Rate
# Quincy University	16.3	32	6880	80
# Bard College	10.4	30	13894	79
# University of Southern California	11.4	10	17007	68

```
testing_data <- split_data$testing_data
head(testing_data, n = 3)
```

#	Private	Apps	Accept	Enroll	Top10perc	Top25perc
# Washington State University	No	6540	5839	2440	31	70
# Saint Francis College IN	Yes	213	166	85	13	36
# St. Mary's College of California	Yes	2643	1611	465	36	80
#	F.Undergrad	P.Undergrad	Outstate	Room.Board		
# Washington State University	14445	1344	8200	4210		
# Saint Francis College IN	513	247	8670	3820		
# St. Mary's College of California	2615	248	13332	6354		
#	Books	Personal	PhD	Terminal	S.F.Ratio	
# Washington State University	800	2719	84	87	16.9	
# Saint Francis College IN	450	1000	43	78	12.5	
# St. Mary's College of California	630	1584	88	89	16.1	
#	perc.alumni	Expend	Grad.Rate			
# Washington State University	30	10912	56			
# Saint Francis College IN	4	7440	48			
# St. Mary's College of California	17	9619	78			

- (b) Fit a linear model using least squares on the training set, and report the test error obtained.

```
formula = Apps ~ .
linear_model <- lm(formula, data = training_data)
calculate_mean_squared_error(linear_model)
```

```
# [1] 1043968
```


```
y_hat_linear <- predict(
  object = linear_model,
  testing_data
)
testing_vector_of_numbers_of_applications <-
  colleges$Apps[vector_of_indices_of_testing_data]
y <- testing_vector_of_numbers_of_applications
y_bar <- mean(testing_vector_of_numbers_of_applications)
R_squared <- 1 - mean((y_hat_linear - y)^2) / mean((y_bar - y)^2)
R_squared
```

```
# [1] 0.8744088
```

The resulting mean squared error of prediction $MSEP_{linear} = 1,043,968$.

- (c) Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```
full_model_matrix <- model.matrix(object = formula, data = colleges)[, -1]
the_cv.glmnet <- glmnet::cv.glmnet(
  x = full_model_matrix,
  y = colleges$Apps,
  alpha = 0
)
plot(the_cv.glmnet)
```



```
optimal_value_of_lambda <- the_cv.glmnet$lambda.min
optimal_value_of_lambda
```

```
# [1] 364.8993
```

```
training_data_frame <- colleges[vector_of_indices_of_training_data, ]
training_model_matrix <- model.matrix(
  object = formula,
  data = training_data_frame
)[, -1]
training_vector_of_numbers_of_applications <-
  colleges$Apps[vector_of_indices_of_training_data]
polynomial_ridge_regression_model <- glmnet::glmnet(
  x = training_model_matrix,
  y = training_vector_of_numbers_of_applications,
```

```

    alpha = 0
  )
  testing_data_frame <- colleges[vector_of_indices_of_testing_data, ]
  testing_model_matrix <- model.matrix(
    object = formula,
    data = testing_data_frame
  )[, -1]
  vector_of_predicted_numbers_of_applications <- predict(
    object = polynomial_ridge_regression_model,
    s = optimal_value_of_lambda,
    newx = testing_model_matrix
  )[, 1]
  vector_of_residuals <-
    vector_of_predicted_numbers_of_applications -
    testing_vector_of_numbers_of_applications
  vector_of_squared_residuals <- vector_of_residuals^2
  mean_squared_error <- mean(vector_of_squared_residuals)
  mean_squared_error

```

```
# [1] 1071002
```

```

y_hat_ridge <- vector_of_predicted_numbers_of_applications
R_squared <- 1 - mean((y_hat_ridge - y)^2) / mean((y_bar - y)^2)
R_squared

```

```
# [1] 0.8908082
```


The resulting mean squared error of prediction $MSEP_{ridge} = 1,071,002$. $MSEP_{linear} < MSEP_{ridge}$.

- (d) Fit a lasso model on the training set, with λ chosen by cross-validation. Report the test error obtained, along with the number of non-zero coefficient estimates.

```

full_model_matrix <- model.matrix(object = formula, data = colleges)[, -1]
the_cv.glmnet <- glmnet::cv.glmnet(
  x = full_model_matrix,
  y = colleges$Apps,
  alpha = 1
)
plot(the_cv.glmnet)

```

```
optimal_value_of_lambda <- the_cv.glmnet$lambda.min
optimal_value_of_lambda
```

```
# [1] 12.51776
```

```
training_model_matrix <- model.matrix(
  object = formula,
  data = training_data_frame
)[, -1]
training_vector_of_numbers_of_applications <-
  colleges$Apps[vector_of_indices_of_training_data]
testing_vector_of_numbers_of_applications <-
  colleges$Apps[vector_of_indices_of_testing_data]
polynomial_lasso_regression_model <- glmnet::glmnet(
  x = training_model_matrix,
  y = training_vector_of_numbers_of_applications,
  alpha = 1
)
testing_model_matrix <- model.matrix(
  object = formula,
  data = testing_data_frame
)[, -1]
vector_of_predicted_numbers_of_applications <- predict(
  object = polynomial_lasso_regression_model,
  s = optimal_value_of_lambda,
  newx = testing_model_matrix
```

```
)[, 1]
vector_of_residuals <-
  vector_of_predicted_numbers_of_applications -
  testing_vector_of_numbers_of_applications
vector_of_squared_residuals <- vector_of_residuals^2
mean_squared_error <- mean(vector_of_squared_residuals)
mean_squared_error
```

```
# [1] 1168351
```

```
predict(
  object = polynomial_lasso_regression_model,
  s = optimal_value_of_lambda,
  type = "coefficients"
)
```

```
# 18 x 1 sparse Matrix of class "dgCMatrix"
#               s1
# (Intercept) -4.772317e+02
# PrivateYes   -4.945991e+02
# Accept       1.550834e+00
# Enroll       -4.885202e-01
# Top10perc    3.994779e+01
# Top25perc    -5.578160e+00
# F.Undergrad  1.402583e-03
# P.Undergrad  4.404712e-02
# Outstate     -7.179337e-02
# Room.Board   1.435632e-01
# Books        6.641890e-02
# Personal     9.897576e-03
# PhD          -6.550167e+00
# Terminal     -3.749746e+00
# S.F.Ratio    4.409518e+00
# perc.alumni  -1.416008e-01
# Expend       6.819184e-02
# Grad.Rate    6.068215e+00
```

```
y_hat_lasso <- vector_of_predicted_numbers_of_applications
R_squared <- 1 - mean((y_hat_lasso - y)^2) / mean((y_bar - y)^2)
R_squared
```

```
# [1] 0.8808833
```

The number of non-zero coefficients $M = 17$, which is equal to the number of non-zero coefficients. The resulting mean squared error of prediction $MSEP_{lasso} = 1,168,351$. $MSEP_{linear} < MSEP_{ridge} < MSEP_{lasso}$.


- (e) Fit a PCR model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
library(pls)
```

```
#
# Attaching package: 'pls'

# The following object is masked from 'package:stats':
#
#   loadings

the_mvr <- pcr(formula, data = training_data_frame, scale = TRUE, validation = "CV")
validationplot(the_mvr, val.type = "MSEP")
```



```
vector_of_predicted_numbers_of_applications <- predict(
  object = the_mvr,
  testing_data_frame,
  ncomp = 17
)
vector_of_residuals <-
  vector_of_predicted_numbers_of_applications -
  testing_vector_of_numbers_of_applications
vector_of_squared_residuals <- vector_of_residuals^2
mean_squared_error <- mean(vector_of_squared_residuals)
mean_squared_error
```

```
# [1] 1231856
```

```

y_hat_PCR <- vector_of_predicted_numbers_of_applications
R_squared <- 1 - mean((y_hat_PCR - y)^2) / mean((y_bar - y)^2)
R_squared

```

```
# [1] 0.8744088
```


The resulting mean squared error of prediction $MSEP_{PCR} = 1,231,856$. $MSEP_{linear} < MSEP_{ridge} < MSEP_{lasso} < MSEP_{PCR}$.

- (f) Fit a PLS model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```

library(pls)
the_mvr <- plsr(formula, data = training_data_frame, scale = TRUE, validation = "CV")
validationplot(the_mvr, val.type = "MSEP")

```



```

vector_of_predicted_numbers_of_applications <- predict(
  object = the_mvr,
  testing_data_frame,
  ncomp = 17
)
vector_of_residuals <-
  vector_of_predicted_numbers_of_applications -
  testing_vector_of_numbers_of_applications
vector_of_squared_residuals <- vector_of_residuals^2
mean_squared_error <- mean(vector_of_squared_residuals)
mean_squared_error

```

```
# [1] 1231856
```

```
y_hat_PLSR <- vector_of_predicted_numbers_of_applications  
R_squared <- 1 - mean((y_hat_PLSR - y)^2) / mean((y_bar - y)^2)  
R_squared
```

```
# [1] 0.8744088
```

The number of non-zero coefficients estimates is 17. The resulting mean squared error of prediction $MSEP_{PLSR} = 1,231,856$. $MSEP_{linear} < MSEP_{ridge} < MSEP_{lasso} < MSEP_{PCR}$; $MSEP_{PCR} = MSEP_{PLSR}$.

- (g) Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these five approaches?

As above, $MSEP_{linear} < MSEP_{ridge} < MSEP_{lasso} < MSEP_{PCR}$; $MSEP_{PCR} = MSEP_{PLSR}$. The mean squared error of prediction of the linear model $MSEP_{linear} = 1,043,968$. The mean squared error of prediction of the PLSR model $MSEP_{PLSR} = 1,231,856$. The rate of difference between these errors is -0.153 . There is a strong effect for each model; $R^2_{linear} = R^2_{PCR} = R^2_{PLSR}$; $R^2_{linear} < R^2_{lasso} < R^2_{ridge}$.