

DS-6030 Homework Module 7

Tom Lever

07/08/2023

DS 6030 | Spring 2023 | University of Virginia

8. In the lab, a classification tree was applied to the Carseats data set after converting Sales into a qualitative response variable.

Now we will seek to predict Sales using regression trees and related approaches, treating the response as a quantitative variable.

- (a) Split the data set into a training set and a test set.


```
set.seed(1)
library(ISLR2)
library(TomLeversRPackage)
training_and_testing_data <- split_data_set_into_training_and_testing_data(
  Carseats,
  proportion_of_training_data = 0.9
)
training_data <- training_and_testing_data$training_data
testing_data <- training_and_testing_data$testing_data
```

- (b) Fit a regression tree to the training set. Plot the tree, and interpret the results. What test MSE do you obtain?

```
library(tree)
full_tree <- tree(Sales ~ ., data = training_data)
summary(full_tree)

#
# Regression tree:
# tree(formula = Sales ~ ., data = training_data)
# Variables actually used in tree construction:
# [1] "ShelveLoc" "Price" "Age" "Income" "CompPrice"
# [6] "Advertising"
# Number of terminal nodes: 17
# Residual mean deviance: 2.653 = 910.1 / 343
# Distribution of residuals:
# Min. 1st Qu. Median Mean 3rd Qu. Max.
# -5.18600 -1.09000 0.05305 0.00000 1.08300 4.63100

plot(full_tree)
text(full_tree, pretty = 0)
```



```
vector_of_predicted_sales <- predict(full_tree, newdata = testing_data)
vector_of_actual_sales <- testing_data$Sales
calculate_mean_squared_error(vector_of_predicted_sales, vector_of_actual_sales)


# [1] 4.896065
```

When shelf location is good and price is less than 109.5 monetary units, our tree predicts that 12.190 thousand child car seats will be sold at each location in each time period.


The test Mean Squared Error of our tree when predicting sales is 4.896 *thousand*².

- (c) Use cross-validation in order to determine the optimal level of tree complexity. Does pruning the tree improve the test MSE?

```
object_of_types_prune_and_tree_sequence <- cv.tree(full_tree)
vector_of_sizes <- object_of_types_prune_and_tree_sequence$size
vector_of_deviances <- object_of_types_prune_and_tree_sequence$dev
plot(vector_of_sizes, vector_of_deviances, type = "l")
index_of_minimum_deviance <- which.min(vector_of_deviances)
optimal_size <-
  vector_of_sizes[index_of_minimum_deviance]
minimum_deviance <- min(vector_of_deviances)
points(
  optimal_size,
  minimum_deviance,
  col = "red",
  cex = 2,
  pch = 20
)
```



```
pruned_tree <- prune.tree(full_tree, best = optimal_size)
plot(pruned_tree)
text(pruned_tree, pretty = 0)
```



```
vector_of_predicted_sales <- predict(pruned_tree, newdata = testing_data)
calculate_mean_squared_error(vector_of_predicted_sales, vector_of_actual_sales)
```

```
# [1] 6.785063
```

The test Mean Squared Error for the pruned tree is greater and less desirable than the Mean Squared Error for the full tree.

- (d) Use the bagging approach in order to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important.

Per *An Introduction to Statistical Learning* (Second Edition), bagging “is simply a special case of a random forest with [the number of variables randomly sampled as candidates at each split] $m = p$ the number of predictors.”

```
library(randomForest)
```

```
# randomForest 4.7-1.1
```

```
# Type rfNews() to see new features/changes/bug fixes.
```

```
index_of_column_Sales <-
  get_index_of_column_of_data_frame(training_data, "Sales")
data_frame_of_predictors <- training_data[, -index_of_column_Sales]
data_frame_of_sales <- training_data[, index_of_column_Sales]
number_of_predictors <- ncol(data_frame_of_predictors)
get_test_MSE_and_vector_of_ordered_percent_increases_in_MSE_for_random_forest <-
  function(mtry) {
    the_randomForest <- randomForest(
      formula = Sales ~ .,
      data = training_data,
      mtry = mtry,
      importance = TRUE
    )
    vector_of_predicted_sales <-
      predict(the_randomForest, newdata = testing_data)
    test_MSE <- calculate_mean_squared_error(
      vector_of_predicted_sales,
      vector_of_actual_sales
    )
    matrix_of_importance_metrics <- importance(the_randomForest)
    vector_of_percent_increases_in_MSE <-
      matrix_of_importance_metrics[, "%IncMSE"]
    vector_of_indices_of_ordered_percent_increases_in_MSE <-
      order(vector_of_percent_increases_in_MSE, decreasing = TRUE)
    vector_of_ordered_percent_increases_in_MSE <-
      vector_of_percent_increases_in_MSE[
        vector_of_indices_of_ordered_percent_increases_in_MSE
      ]
    list_of_test_MSE_and_vector_of_ordered_percent_increases_in_MSE_for_random_forest <-
      list(
        test_MSE = test_MSE,
        vector_of_ordered_percent_increases_in_MSE =
          vector_of_ordered_percent_increases_in_MSE
      )
    return(
```

```

        list_of_test_MSE_and_vector_of_ordered_percent_in_MSE_for_random_forest
    )
}
get_test_MSE_and_vector_of_ordered_percent_in_MSE_for_random_forest(
    mtry = number_of_predictors
)

```

```

# $test_MSE
# [1] 2.912954
#
# $vector_of_ordered_percent_in_MSE
#   ShelfLoc      Price  CompPrice Advertising      Age      Income
# 81.267672  79.323197  38.593171  25.822124  25.716498  14.313945
#   Education      US      Urban  Population
#   2.885808   2.270115  -1.691179  -2.211183

```

The test Mean Squared Error for our bootstrap aggregation (BAG) is 2.913, which is 0.595 of the MSE for our full tree and 0.429 of the MSE for our pruned tree.

According to [In a random forest, is larger %IncMSE better or worse?](#), “%IncMSE is the most robust and informative measure. IT is the increase in mse of predictions(estimated with out-of-bag-CV) as a result of variable j being permuted(values randomly shuffled)... the higher the number, the more important.”

%IncMSE is highest for *ShelveLoc* followed by *Price*; *ShelveLoc* and *Price* are the two most important variables.

- (e) Use random forests to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important. Describe the effect of `m`, the number of variables considered at each split, on the error rate obtained.

```

data_frame_of_values_of_mtry_and_test_MSEs <- data.frame(
    matrix(NA, nrow = number_of_predictors, ncol = 2)
)
colnames(data_frame_of_values_of_mtry_and_test_MSEs) <- c("mtry", "test_MSE")
for (mtry in 1:number_of_predictors) {
    print(paste("mtry: ", mtry, sep = ""))
    data_frame_of_values_of_mtry_and_test_MSEs[mtry, "mtry"] <- mtry
    test_MSE_and_vector_of_ordered_percent_in_MSE <-
        get_test_MSE_and_vector_of_ordered_percent_in_MSE_for_random_forest(
            mtry = mtry
        )
    test_MSE <- test_MSE_and_vector_of_ordered_percent_in_MSE$test_MSE
    vector_of_ordered_percent_in_MSE <-
        test_MSE_and_vector_of_ordered_percent_in_MSE$
            vector_of_ordered_percent_in_MSE
    print(vector_of_ordered_percent_in_MSE)
    data_frame_of_values_of_mtry_and_test_MSEs[mtry, "test_MSE"] <- test_MSE
}

```

```

# [1] "mtry: 1"
#   ShelfLoc      Price      Age Advertising  CompPrice      US
# 27.4145011 22.6376857 12.1022390 11.8362136  9.1428466  6.3777495
#   Income  Education      Urban  Population
# 5.7156306  2.2635817 -0.2072851 -0.9091637
# [1] "mtry: 2"
#   ShelfLoc      Price Advertising      Age  CompPrice      Income

```

```

# 44.880078 37.860838 17.212010 16.137100 14.665327 6.413410
# US Education Urban Population
# 6.210358 1.257871 -1.022202 -1.798896
# [1] "mtry: 3"
# ShelfLoc Price Advertising Age CompPrice Income
# 57.917795 48.708554 19.677702 19.528156 18.673141 7.957699
# US Education Population Urban
# 5.868734 3.956120 -2.433051 -2.642151
# [1] "mtry: 4"
# ShelfLoc Price CompPrice Age Advertising Income
# 60.737660 56.391682 23.604401 21.854985 18.236436 8.220434
# US Education Urban Population
# 6.106552 1.320007 -1.574042 -1.839350
# [1] "mtry: 5"
# ShelfLoc Price CompPrice Age Advertising Income
# 68.562969 61.860220 25.200712 23.425230 20.589795 10.621514
# US Education Population Urban
# 5.444433 3.662041 -1.106284 -2.907471
# [1] "mtry: 6"
# ShelfLoc Price CompPrice Age Advertising Income
# 76.681110 67.104963 29.750847 24.706524 21.958387 11.035165
# US Education Urban Population
# 3.985181 2.732473 -1.483772 -2.223225
# [1] "mtry: 7"
# ShelfLoc Price CompPrice Age Advertising Income
# 79.7220186 72.1876642 34.2439225 24.6451061 23.4062353 11.4780558
# US Education Population Urban
# 4.3765943 2.7168918 -0.4954921 -1.9197438
# [1] "mtry: 8"
# ShelfLoc Price CompPrice Age Advertising Income
# 79.5222605 74.7577736 33.7322095 24.6121401 24.4295814 14.3477921
# US Education Population Urban
# 4.9868483 2.4070209 -0.2734359 -1.8060475
# [1] "mtry: 9"
# ShelfLoc Price CompPrice Age Advertising Income
# 83.989493 78.130492 38.157897 26.980376 23.366049 13.876704
# Education US Population Urban
# 3.103712 2.183155 -1.673313 -2.381346
# [1] "mtry: 10"
# ShelfLoc Price CompPrice Advertising Age Income
# 81.363209 77.647089 38.547174 27.190982 22.778387 14.981041
# US Education Urban Population
# 4.123539 3.332881 -1.602569 -2.302125

```

```
print(data_frame_of_values_of_mtry_and_test_MSEs)
```

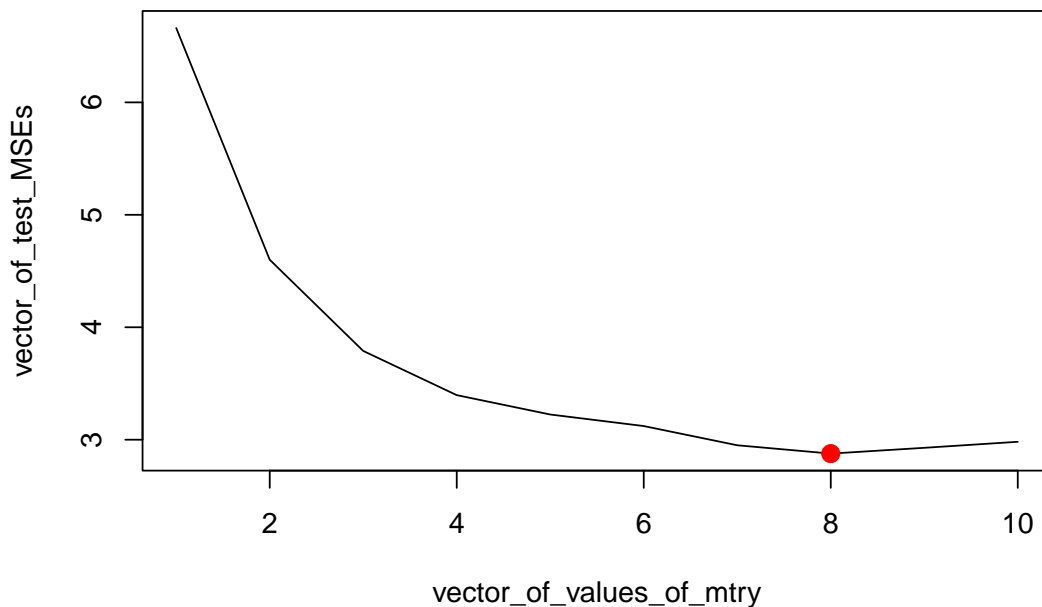
```

# mtry test_MSE
# 1 1 6.661409
# 2 2 4.600594
# 3 3 3.789455
# 4 4 3.396240
# 5 5 3.224155
# 6 6 3.121334
# 7 7 2.950383

```

```
# 8      8 2.876444
# 9      9 2.927813
# 10     10 2.981307
```

```
vector_of_values_of_mtry <- data_frame_of_values_of_mtry_and_test_MSEs$mtry
vector_of_test_MSEs <- data_frame_of_values_of_mtry_and_test_MSEs$test_MSE
plot(
  x = vector_of_values_of_mtry,
  y = vector_of_test_MSEs,
  type = "l"
)
index_of_minimum_test_MSE <- which.min(vector_of_test_MSEs)
optimal_value_of_mtry <-
  vector_of_values_of_mtry[index_of_minimum_test_MSE]
minimum_test_MSE <- min(vector_of_test_MSEs)
points(
  optimal_value_of_mtry,
  minimum_test_MSE,
  col = "red",
  cex = 2,
  pch = 20
)
```



See above plot for test Mean Squared Errors for different values of the number of variables randomly sampled as candidates at each split m . Test MSE decreases parabolically with number of variables to a minimum for $m = 8$. In all cases *ShelveLoc* and *Price* are the most important predictors.

(f) Now analyze the data using BART, and report your results. (skip this exercise)

9. This problem involves the OJ data set which is part of the ISLR package.

- (a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
training_and_testing_data <- split_data_set_into_training_and_testing_data(
  OJ,
  number_of_training_data = 800
)
training_data <- training_and_testing_data$training_data
testing_data <- training_and_testing_data$testing_data
```

- (b) Fit a tree to the training data, with *Purchase* as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics about the tree, and describe the results obtained. What is the training error rate? How many terminal nodes does the tree have?

```
full_tree <- tree(Purchase ~ ., data = training_data)
summary(full_tree)
```

```
#
# Classification tree:
# tree(formula = Purchase ~ ., data = training_data)
# Variables actually used in tree construction:
# [1] "LoyalCH"      "SalePriceMM" "SpecialCH"    "PriceDiff"    "STORE"
# Number of terminal nodes: 9
# Residual mean deviance: 0.7139 = 564.7 / 791
# Misclassification error rate: 0.1575 = 126 / 800
```

Our full tree is a classification tree that predicts whether a customer will purchase Citrus Hill or Minute Maid orange juice. A tree is grown by binary recursive partitioning using the response in the specified formula, *Purchase*, and choosing splits from the terms of the right-hand-side, which in our case are all terms besides *Purchase*. The predictors actually used in tree construction are *LoyalCH*, *SalePriceMM*, *SpecialCH*, *PriceDiff*, and *STORE*. *Purchase* is a factor with levels *CH* and *MM* indicating whether the customer purchased Citrus Hill or Minute Maid Orange Juice. *LoyalCH* seems to be a rate of customer brand loyalty for CH between 0 and 1. *SalePriceMM* seems to be the net sale price of Minute Maid orange juice in dollars. *SpecialCH* seems to be an indicator of whether or not there is a special on Citrus Hill orange juice. *PriceDiff* seems to be net sale price of Minute Maid orange juice less net sale price of Citrus Hill orange juice in dollars. *STORE* seems to be a categorical variable indicating at which of 5 possible stores the purchase occurred. In our full tree there are 9 terminal nodes / leaves. The deviance of our full tree is 564.7. A small deviance indicates a tree that provides a good fit to the training data. The residual mean deviance for our full tree is $564.7 / (800 - 9)$. The training error rate / misclassification error rate for our full tree is $126 / 800$.

- (c) Type in the name of the tree object in order to get a detailed text output. Pick one of the terminal nodes, and interpret the information displayed.

```
full_tree

# node), split, n, deviance, yval, (yprob)
#      * denotes terminal node
#
# 1) root 800 1077.00 CH ( 0.60000 0.40000 )
#   2) LoyalCH < 0.5036 341 371.50 MM ( 0.23460 0.76540 )
#     4) LoyalCH < 0.282272 167 114.20 MM ( 0.10778 0.89222 ) *
#     5) LoyalCH > 0.282272 174 226.60 MM ( 0.35632 0.64368 )
#       10) SalePriceMM < 2.04 97 101.40 MM ( 0.21649 0.78351 )
#         20) SpecialCH < 0.5 75 58.90 MM ( 0.13333 0.86667 ) *
```




```
#      21) SpecialCH > 0.5 22   30.50 CH ( 0.50000 0.50000 ) *
#      11) SalePriceMM > 2.04 77 106.40 CH ( 0.53247 0.46753 ) *
#      3) LoyalCH > 0.5036 459 352.10 CH ( 0.87146 0.12854 )
#      6) PriceDiff < -0.39 22   27.52 MM ( 0.31818 0.68182 ) *
#      7) PriceDiff > -0.39 437 285.40 CH ( 0.89931 0.10069 )
#      14) LoyalCH < 0.764572 175 172.30 CH ( 0.80571 0.19429 )
#      28) SalePriceMM < 2.125 106 126.30 CH ( 0.71698 0.28302 ) *
#      29) SalePriceMM > 2.125 69   30.55 CH ( 0.94203 0.05797 ) *
#      15) LoyalCH > 0.764572 262 84.93 CH ( 0.96183 0.03817 )
#      30) STORE < 1.5 133    0.00 CH ( 1.00000 0.00000 ) *
#      31) STORE > 1.5 129   70.35 CH ( 0.92248 0.07752 ) *
```

A terminal node / leaf / branch to leaf is indicated by an asterisk. Because `full_tree` outputs information for each of two branches for each internal node / node other than the root node / trunk and the leaves, we speak in terms of branches. Let us consider the terminal node / leaf / branch to leaf labeled 4. The prediction of the full tree associated with this branch to leaf is *MM*. We arrive at this branch when the split criterion *LoyalCH* is less than 0.504 and less than 0.282. The split criterion associated with this branch is *LoyalCH* < 0.282. The number of observations / purchases in our training data set is 800. Of those purchases, 341 purchases are by customers with loyalty to Citrus Hill less than 0.504. Of those purchases, 167 purchases are by customers with loyalty to Citrus Hill less than 0.282. The number of purchases associated with our branch is 167 with a deviance of 114.2. 0.108 of purchases associated with our branch were of Citrus Hill orange juice. 0.892 of purchases associated with our branch were of Minute Maid orange juice.

(d) Create a plot of the tree, and interpret the results.

```
plot(full_tree)
text(full_tree, pretty = 0)
```



Per our full tree, the most important predictor of whether a customer will purchase Citrus Hill

or Minute Maid orange juice is loyalty to Citrus Hill. The split criterion for the first non-root / internal node / the first pair of branches is *LoyalCH*. The split criteria for the second internal node in the second echelon is also *LoyalCH*.

- (e) Predict the response on the test data, and produce a confusion matrix comparing the test labels to the predicted test labels. What is the test error rate?

```
vector_of_predicted_purchases <- predict(full_tree, testing_data, type = "class")
vector_of_actual_purchases <- testing_data$Purchase
confusion_matrix <- table(vector_of_predicted_purchases, vector_of_actual_purchases)
confusion_matrix
```

```
#               vector_of_actual_purchases
# vector_of_predicted_purchases  CH  MM
#               CH 144  27
#               MM  29  70
```

```
number_of_purchases_of_Citrus_Hill_orange_juice_predicted_correctly <- 144
number_of_purchases_of_Minute_Maid_orange_juice_predicted_correctly <- 70
number_of_purchases_predicted_correctly <-
  number_of_purchases_of_Citrus_Hill_orange_juice_predicted_correctly +
  number_of_purchases_of_Minute_Maid_orange_juice_predicted_correctly
number_of_purchases <- nrow(testing_data)
test_accuracy <- number_of_purchases_predicted_correctly / number_of_purchases
test_error_rate <- 1 - test_accuracy
test_error_rate
```

```
# [1] 0.2074074
```

The test error rate is about 0.207.

- (f) Apply the `cv.tree()` function to the training set in order to determine the optimal tree size.

```
object_of_types_prune_and_tree_sequence <-
  cv.tree(full_tree, FUN = prune.misclass)
vector_of_sizes <- object_of_types_prune_and_tree_sequence$size
vector_of_numbers_of_errors <- object_of_types_prune_and_tree_sequence$dev
minimum_number_of_errors <- min(vector_of_numbers_of_errors)
index_of_minimum_number_of_errors <- which.min(vector_of_numbers_of_errors)
optimal_size <-
  vector_of_sizes[index_of_minimum_number_of_errors]
optimal_size
```

```
# [1] 3
```

The optimal tree size is 3.


- (g) Produce a plot with tree size on the x-axis and cross-validated classification error rate on the y-axis.

```
plot(
  vector_of_sizes,
  vector_of_numbers_of_errors,
  type = "l",
  xlab = "tree size",
  ylab = "cross-validated number of classification errors"
)
points(
```

```

    optimal_size,
    minimum_number_of_errors,
    col = "red",
    cex = 2,
    pch = 20
)

```




- (h) Which tree size corresponds to the lowest cross-validated classification error rate?
 A tree size of 3 corresponds to the lowest cross-validated classification number of errors.
- (i) Produce a pruned tree corresponding to the optimal tree size obtained using cross-validation. If cross-validation does not lead to selection of a pruned tree, then create a pruned tree with five terminal nodes.

```

pruned_tree <- prune.tree(full_tree, best = optimal_size)
plot(pruned_tree)
text(pruned_tree, pretty = 0)

```



- (j) Compare the training error rates between the pruned and unpruned trees. Which is higher?
- (k) Compare the test error rates between the pruned and unpruned trees. Which is higher?