

DS-6030 Homework Module 8

Tom Lever

07/16/2023

DS 6030 | Spring 2023 | University of Virginia

7. In the lab, we applied random forests to the Boston data using `mtry = 6` and using `ntree = 25` and `ntree = 500`.

Create a plot displaying the test error resulting from random forests on this data set for a more comprehensive range of values for `mtry` and `ntree`. You can model your plot after Figure 8.10. Describe the results obtained.

```
library(ISLR2)
library(randomForest)
```

```
# randomForest 4.7-1.1
```

```
# Type rfNews() to see new features/changes/bug fixes.
```

```
library(TomLeversRPackage)
```

```
set.seed(1)
training_and_testing_data <- split_data_set_into_training_and_testing_data(
  Boston,
  proportion_of_training_data = 0.9
)
training_data <- training_and_testing_data$training_data
testing_data <- training_and_testing_data$testing_data
head(training_data, n = 3)
```


```
#      crim zn indus chas   nox   rm  age   dis rad tax ptratio lstat medv
# 505 0.10959 0 11.93    0 0.573 6.794 89.3 2.3889  1 273    21.0  6.48 22.0
# 324 0.28392 0  7.38    0 0.493 5.708 74.3 4.7211  5 287    19.6 11.74 18.5
# 167 2.01019 0 19.58    0 0.605 7.929 96.2 2.0459  5 403    14.7  3.70 50.0
```

```
index_of_column_medv <- get_index_of_column_of_data_frame(training_data, "medv")
data_frame_of_training_predictors <- training_data[, -index_of_column_medv]
number_of_predictors <- ncol(data_frame_of_training_predictors)
data_frame_of_training_response_values <- training_data[, index_of_column_medv]
data_frame_of_testing_predictors <- testing_data[, -index_of_column_medv]
data_frame_of_testing_response_values <- testing_data[, index_of_column_medv]
randomForest_for_mtry_equal_to_number_of_predictors <- randomForest(
  x = data_frame_of_training_predictors,
```

```

    y = data_frame_of_training_response_values,
    xtest = data_frame_of_testing_predictors,
    ytest = data_frame_of_testing_response_values,
    mtry = number_of_predictors,
    ntree = 500
  )
  randomForest_for_mtry_equal_to_half_number_of_predictors <- randomForest(
    x = data_frame_of_training_predictors,
    y = data_frame_of_training_response_values,
    xtest = data_frame_of_testing_predictors,
    ytest = data_frame_of_testing_response_values,
    mtry = number_of_predictors / 2,
    ntree = 500
  )
  randomForest_for_mtry_equal_to_square_root_of_number_of_predictors <- randomForest(
    x = data_frame_of_training_predictors,
    y = data_frame_of_training_response_values,
    xtest = data_frame_of_testing_predictors,
    ytest = data_frame_of_testing_response_values,
    mtry = sqrt(number_of_predictors),
    ntree = 500
  )
  plot(
    x = 1:500,
    y = randomForest_for_mtry_equal_to_number_of_predictors$test$mse,
    ylim = c(10, 20),
    col = "red",
    type = "l",
    xlab = "Number Of Trees",
    ylab = "Test Mean Squared Error",
    main = "Test Mean Squared Error vs. Number Of Trees\nFor Random Forests With Different mtry"
  )
  lines(
    x = 1:500,
    y = randomForest_for_mtry_equal_to_half_number_of_predictors$test$mse,
    col = "green"
  )
  lines(
    x = 1:500,
    y = randomForest_for_mtry_equal_to_square_root_of_number_of_predictors$test$mse,
    col = "blue"
  )
  legend(
    x = "topright",
    legend = c("m = p", "m = p/2", "m = sqrt(p)"),
    col = c("red", "green", "blue"),
    lty = 1
  )
)

```



Above is a plot of Test Mean Squared Error for random forests predicting median value of owner-occupied homes in thousands of dollars based on the other variables of data set `ISLR2::Boston`. Variable `mtry` represents the number of variables considered at each split. Red, green, and blue curves correspond to random forests with `mtry` equal to the number of predictors, half the number of predictors, and the square root of the number of predictors, respectively. For each curve, Test Mean Squared Error decreases exponentially with number of trees. A random forest with `mtry` equal to the number of predictors and a number of trees less than 25 has the lowest Test Mean Squared Error and performs best.

8. This question uses the `Caravan` data set.

- (a) Create a training set consisting of the first 1,000 observations, and a test set consisting of the remaining observations.

```
set.seed(1)
Caravan$Purchase <- ifelse(Caravan$Purchase == "Yes", 1, 0)
training_data <- Caravan[1:1000, ]
testing_data <- Caravan[-c(1:1000), ]
```

- (b) Fit a boosting model to the training set with `Purchase` as the response and the other variables as predictors. Use 1,000 trees, and a shrinkage value of 0.01. Which predictors appear to be the most important?

```
the_gbm.object <- gbm::gbm(
  formula = Purchase ~ .,
  distribution = "adaboost",
  data = training_data,
  n.trees = 1000,
```

```


shrinkage = 0.01
)

# Warning in gbm.fit(x = x, y = y, offset = offset, distribution = distribution,
# : variable 50: PVRAAUT has no variation.

# Warning in gbm.fit(x = x, y = y, offset = offset, distribution = distribution,
# : variable 71: AVRAAUT has no variation.

summary(the_gbm.object)

```



```

#           var      rel.inf
# PPERSAUT PPERSAUT 16.27175732
# MKOOPKLA MKOOPKLA  7.13694190
# MBERMIDD MBERMIDD  6.42793466
# MOPLHOOG MOPLHOOG  5.75847398
# PBRAND    PBRAND   5.23437327
# MGODGE    MGODGE   4.56525360
# MINK3045  MINK3045  4.30566970
# MINKM30   MINKM30   3.84523657
# MAUT1     MAUT1    3.84089359
# MSKC      MSKC     3.61528574
# MBERARBG  MBERARBG  2.93289220
# MAUT2     MAUT2    2.93161159
# MOSTYPE   MOSTYPE  2.58027508
# MSKA      MSKA     2.29951574
# PWAPART   PWAPART  2.20956399
# MBERHOOG  MBERHOOG  2.00974939
# MRELGE    MRELGE   1.88421014

```

#	MGODOV	MGODOV	1.58363134
#	MINKGEM	MINKGEM	1.55185993
#	PBYSTAND	PBYSTAND	1.49720553
#	MGODPR	MGODPR	1.30330440
#	ABRAND	ABRAND	1.30150833
#	MZFONDS	MZFONDS	1.10115776
#	PMOTSCO	PMOTSCO	1.08936547
#	MSKD	MSKD	1.06835520
#	MHHUUR	MHHUUR	1.04244317
#	MSKB1	MSKB1	1.00244172
#	MSKB2	MSKB2	0.81205672
#	MBERBOER	MBERBOER	0.79494102
#	MAUTO	MAUTO	0.76986818
#	MINK4575	MINK4575	0.73016367
#	MRELOV	MRELOV	0.63206423
#	MOPLMIDD	MOPLMIDD	0.62520980
#	MHKOOP	MHKOOP	0.53804087
#	MOSHOOFD	MOSHOOFD	0.52356550
#	MGEMOMV	MGEMOMV	0.44888986
#	MFWEKIND	MFWEKIND	0.44018334
#	MGODRK	MGODRK	0.43872282
#	MRELSA	MRELSA	0.39937493
#	MGEMLEEF	MGEMLEEF	0.37354662
#	MFGEKIND	MFGEKIND	0.35808623
#	MFALLEEN	MFALLEEN	0.33030603
#	APERSAUT	APERSAUT	0.32944210
#	MINK7512	MINK7512	0.31701634
#	MBERARBO	MBERARBO	0.23420446
#	MOPLLAAG	MOPLLAAG	0.18394434
#	MINK123M	MINK123M	0.14264058
#	MBERZELF	MBERZELF	0.11140151
#	MZPART	MZPART	0.07541953
#	MAANTHUI	MAANTHUI	0.00000000
#	PWABEDR	PWABEDR	0.00000000
#	PWALAND	PWALAND	0.00000000
#	PBESAUT	PBESAUT	0.00000000
#	PVRAAUT	PVRAAUT	0.00000000
#	PAANHANG	PAANHANG	0.00000000
#	PTRACTOR	PTRACTOR	0.00000000
#	PWERKT	PWERKT	0.00000000
#	PBROM	PBROM	0.00000000
#	PLEVEN	PLEVEN	0.00000000
#	PPERSONG	PPERSONG	0.00000000
#	PGEZONG	PGEZONG	0.00000000
#	PWAOREG	PWAOREG	0.00000000
#	PZEILPL	PZEILPL	0.00000000
#	PPLEZIER	PPLEZIER	0.00000000
#	PFIETS	PFIETS	0.00000000
#	PINBOED	PINBOED	0.00000000
#	AWAPART	AWAPART	0.00000000
#	AWABEDR	AWABEDR	0.00000000
#	AWALAND	AWALAND	0.00000000
#	ABESAUT	ABESAUT	0.00000000
#	AMOTSCO	AMOTSCO	0.00000000

```
# AVRAAUT    AVRAAUT    0.00000000
# AAANHANG   AAANHANG   0.00000000
# ATTRACTOR  ATTRACTOR  0.00000000
# AWERKT     AWERKT     0.00000000
# ABROM      ABROM      0.00000000
# ALEVEN     ALEVEN     0.00000000
# APERSONG   APERSONG   0.00000000
# AGEZONG    AGEZONG    0.00000000
# AWAOREG    AWAOREG    0.00000000
# AZEILPL    AZEILPL    0.00000000
# APLEZIER   APLEZIER   0.00000000
# AFIETS     AFIETS     0.00000000
# AINBOED    AINBOED    0.00000000
# ABYSTAND   ABYSTAND   0.00000000
```

According to [Understanding Gradient Boosting Machines](#), an “important feature in the gbm modelling is the Variable Importance. Applying the summary function to a gbm output produces both a Variable Importance Table and a Plot of the model. This table [above] ranks the individual variables based on their relative influence, which is a measure indicating the relative importance of each variable in training the model.”

PPERSAUT and *MKOOKLA* appear to our most important predictors.

- (c) Use the boosting model to predict the response on the test data. Predict that a person will make a purchase if the estimated probability of purchase is greater than 20 %. Form a confusion matrix. What fraction of the people predicted to make a purchase do in fact make one? How does this compare with the results obtained from applying KNN or logistic regression to this data set?

```
vector_of_predicted_probabilities <- predict(
  object = the_gbm.object,
  testing_data,
  n.trees = 1000,
  type = "response"
)
vector_of_predictions <- ifelse(
  test = vector_of_predicted_probabilities > 0.2,
  yes = 1,
  no = 0
)
table(testing_data$Purchase, vector_of_predictions)
```

```
#   vector_of_predictions
#      0      1
# 0 4470   63
# 1  270   19
```

The fraction of the people predicted to make a purchase that do in fact make one and the precision of our Generalized Boosting Model for an AdaBoost distribution and threshold of 0.2 $PPV = \frac{TP}{TP+FP} = \frac{19}{19+63} = 0.232$.

```
logistic_regression_model <- glm(
  formula = Purchase ~ .,
  data = training_data,
  family = "binomial"
)
```

```
# Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

vector_of_predicted_probabilities <- predict(
  object = logistic_regression_model,
  testing_data,
  type = "response"
)

# Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
# prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases

vector_of_predictions <- ifelse(
  test = vector_of_predicted_probabilities > 0.2,
  yes = 1,
  no = 0
)
table(testing_data$Purchase, vector_of_predictions)

#      vector_of_predictions
#           0           1
# 0 4183  350
# 1  231   58

```

The precision of a logistic regression model with threshold 0.2 is $PPV = \frac{58}{58+350} = 0.142$. The rate of difference in precision between our Generalized Boosting Model and our logistic regression model is $\frac{0.232-0.142}{0.142} = 0.634$.